



# Introduction to OSG

**Tim Cartwright**

*University of Wisconsin–Madison*

*OSG School Director*

*OSG Campus Coordinator*



So far, we have seen how to use HTC on one cluster

***Don't let computing be a barrier to research!***

*(Sometimes, one cluster is not enough)*

Today, we see what it takes to get more capacity \*

**\* *Caveat:*** *I will focus on compute capacity; Wed. will focus on data.*



# What Researchers Want



Submit locally, run globally





# Yesterday



## Access Point

```
$ ls  
science.sub  
input.dat  
$ condor_q
```




PATH







# The Goal



<https://sweetclipart.com/>  
<https://www.flaticon.com/free-icons/server>



# What That Takes



- ✓ Lots of networked computing capacity
  - OSPool capacity are contributed by campuses!
  - (PATH Facility, used yesterday, is owned by PATH)
- ✓ Capacity owners who are motivated to share
- ✓ Trust among owners, researchers, OSG staff
- ✓ Automation to make it work at scale
- ✓ No allocations or charges would be nice!



# Demo, Part I



# Behind the Curtain





# Reasons for Continuing



- So why learn more about how the OSPool works?
  - May change how you plan to run computing there
  - May change the way you use the Access Point
  - May change how you handle issues that arise
- What is there to learn? (outline of rest of talk)
  - Concepts of OSG, Pool, and Access Point
  - How the OSPool gets resources
  - How the OSPool differs from a local cluster



# What is OSG?





# OSG Defined, Version 1



- OSG Consortium – in this view, OSG *is* people:
  - Users: individuals PI/students to collabs. of 1,000s
  - Resource owners/contributors
  - Team: provides infrastructure, support, features, ...







# OSG Defined, Version 2



- Pools of capacity
  - **Capacity:** compute, storage, and other systems that can be used for research workflows
  - **Services:** software infrastructure that manages capacity and makes features available



<https://www.pngall.com/wp-content/uploads/5/Server-Rack-PNG-Free-Image.png>



# OSG Defined, Version 3



- OSG Access Point
  - Where you go to do computing
  - Has access to capacity (constantly changing)
  - Provides means for accessing data (see Wednesday)

```
[tim.cartwright@ap40 ~]$ condor_version  
$CondorVersion: 10.7.0 2023-07-10 BuildID: 659788 PackageID: 10.7.0-0.659788 RC $  
$CondorPlatform: x86_64_AlmaLinux8 $  
[tim.cartwright@ap40 ~]$ █
```

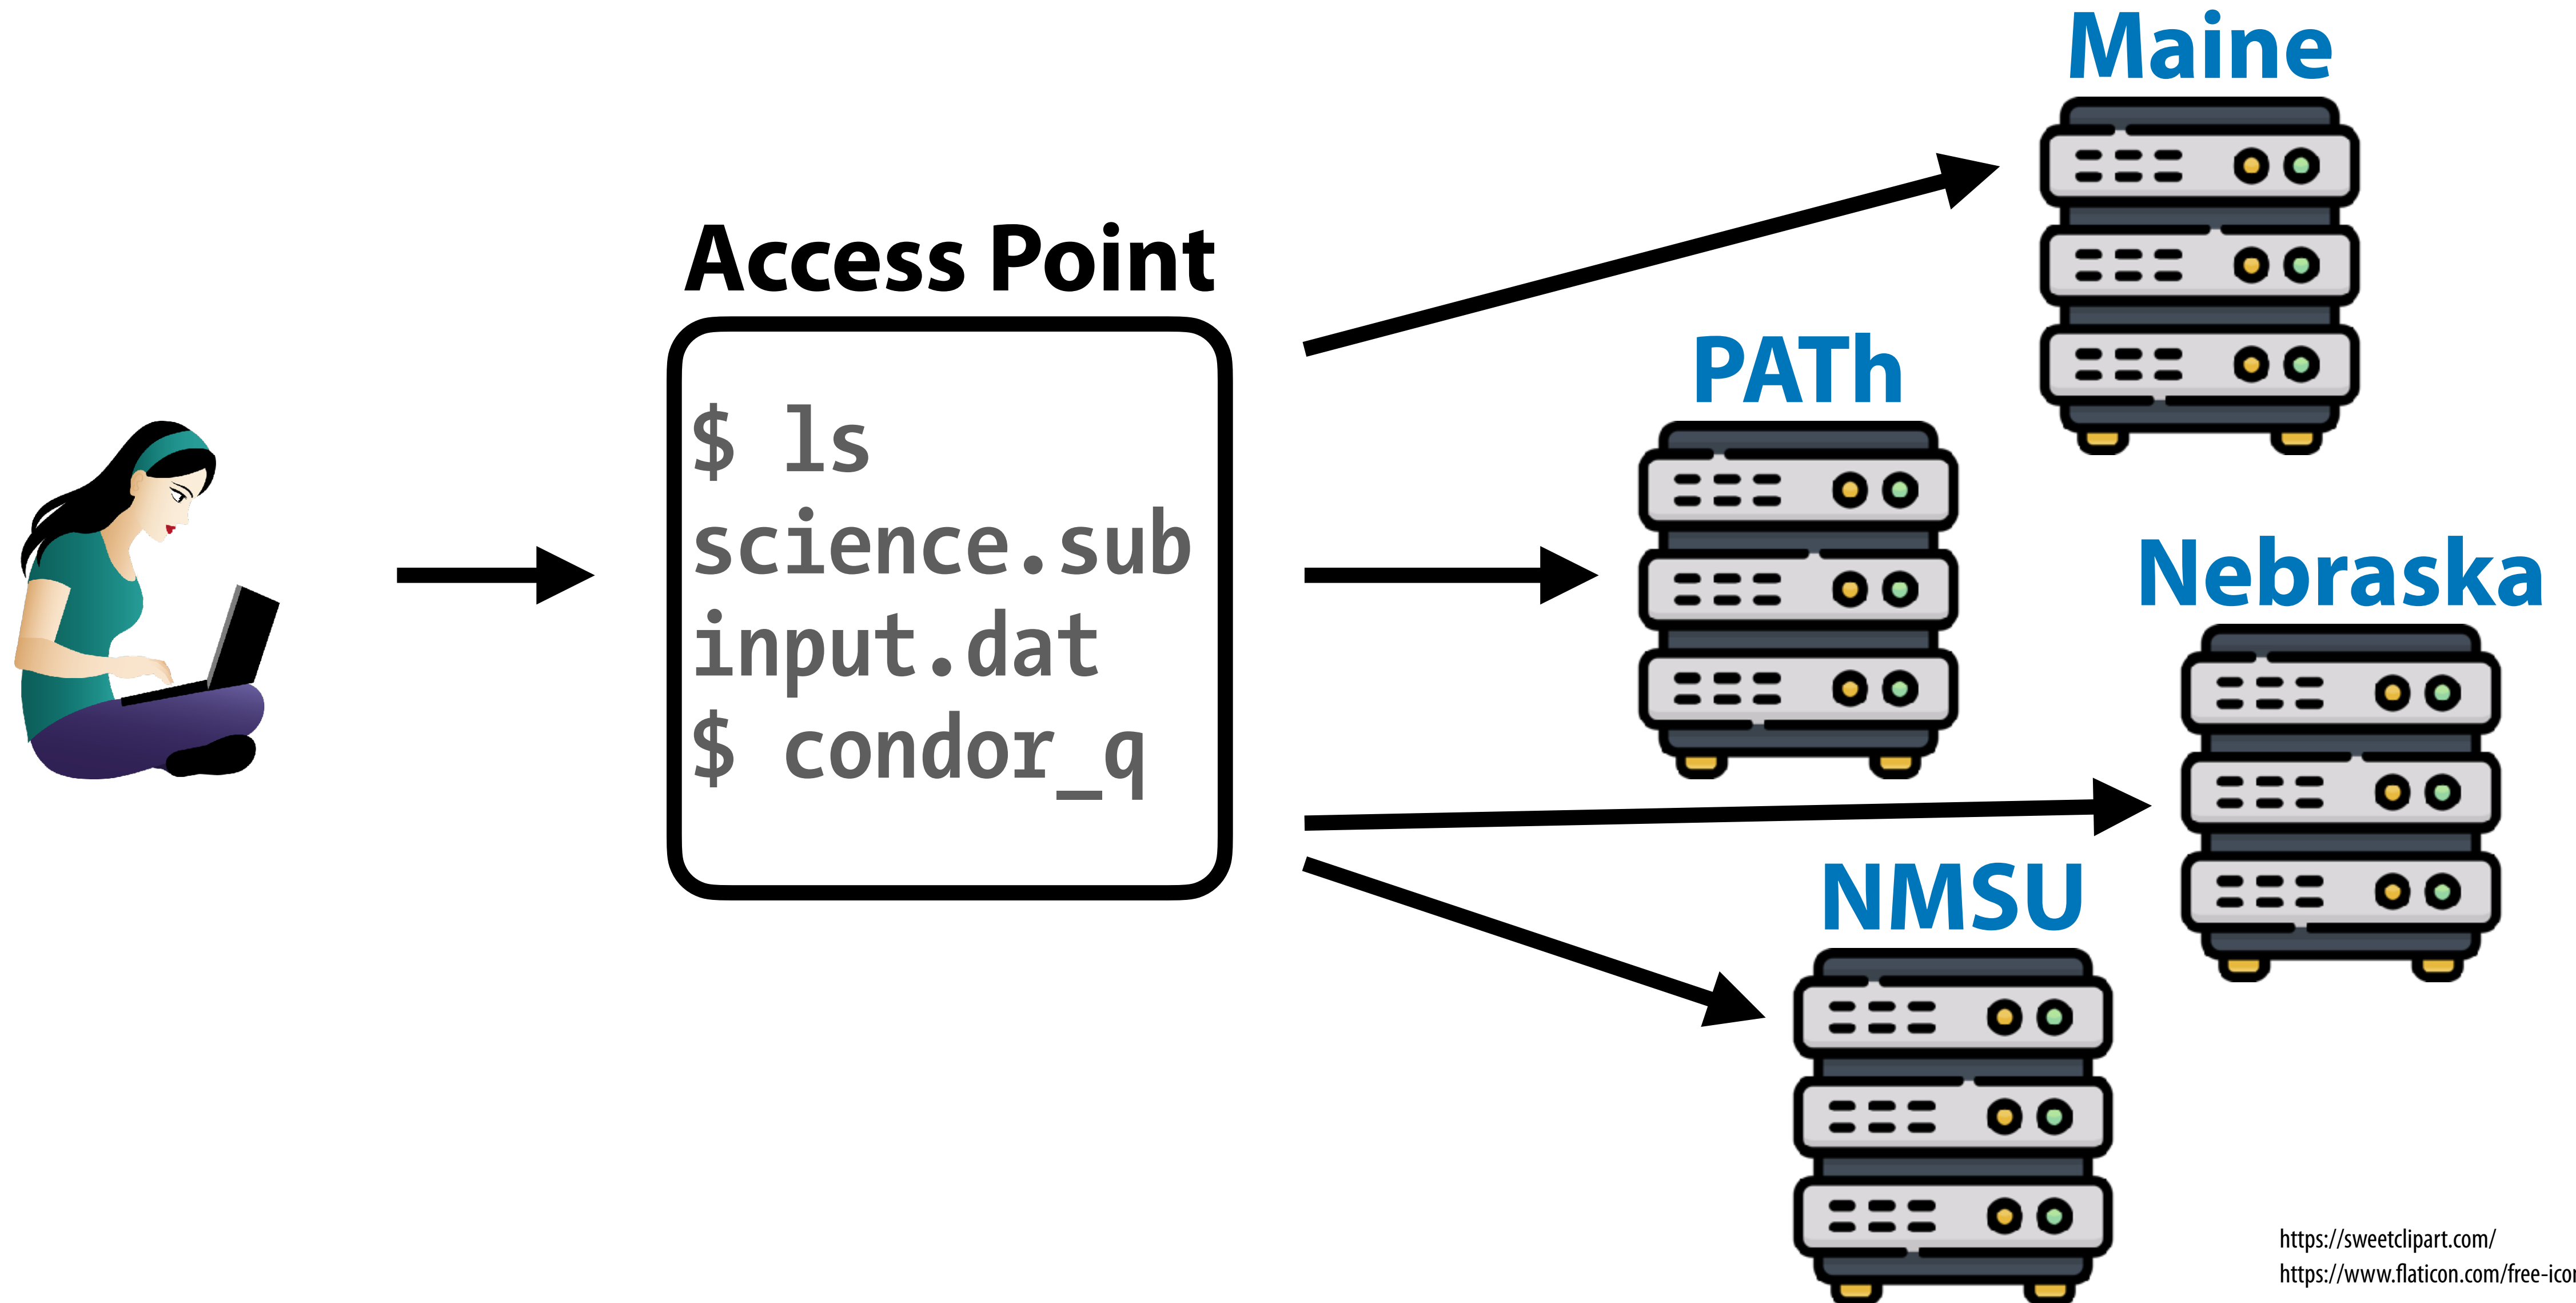




# Getting Capacity for OSPool



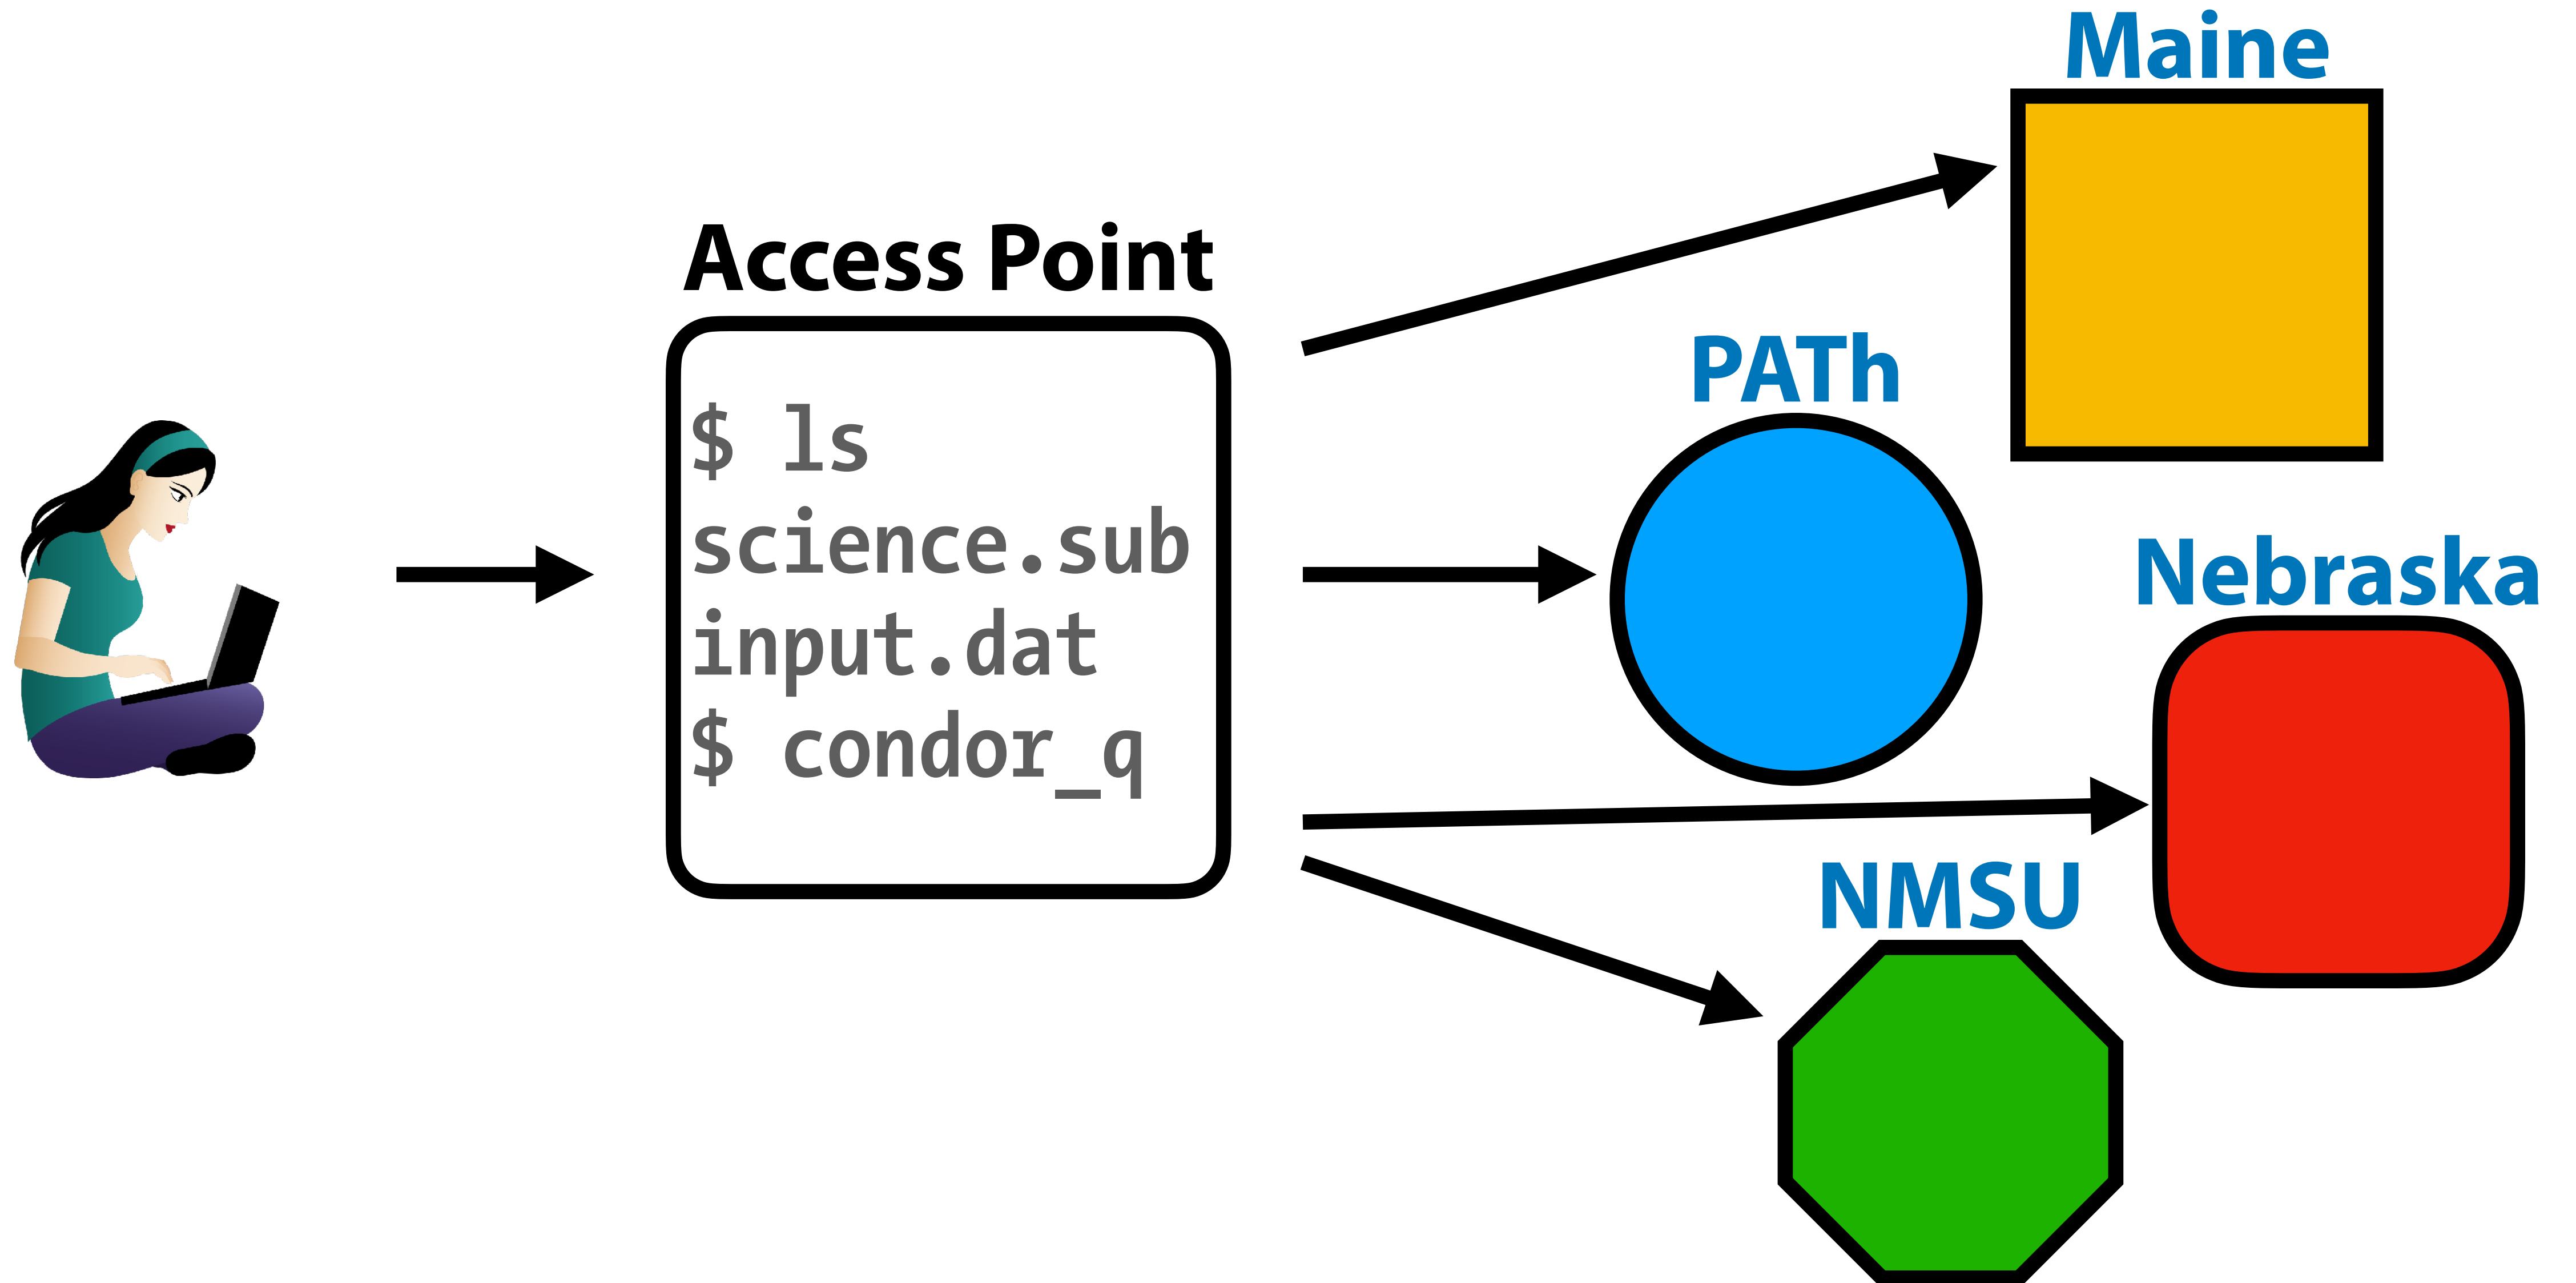
# The Goal



<https://sweetclipart.com/>  
<https://www.flaticon.com/free-icons/server>



# The Goal





## 1. Before OSPool

Nothing available at Wisc. 😭

### Access Point

Job1.0  
Job1.1  
Job1.2  
Job1.3  
...  
Job1.1999

### Wisconsin

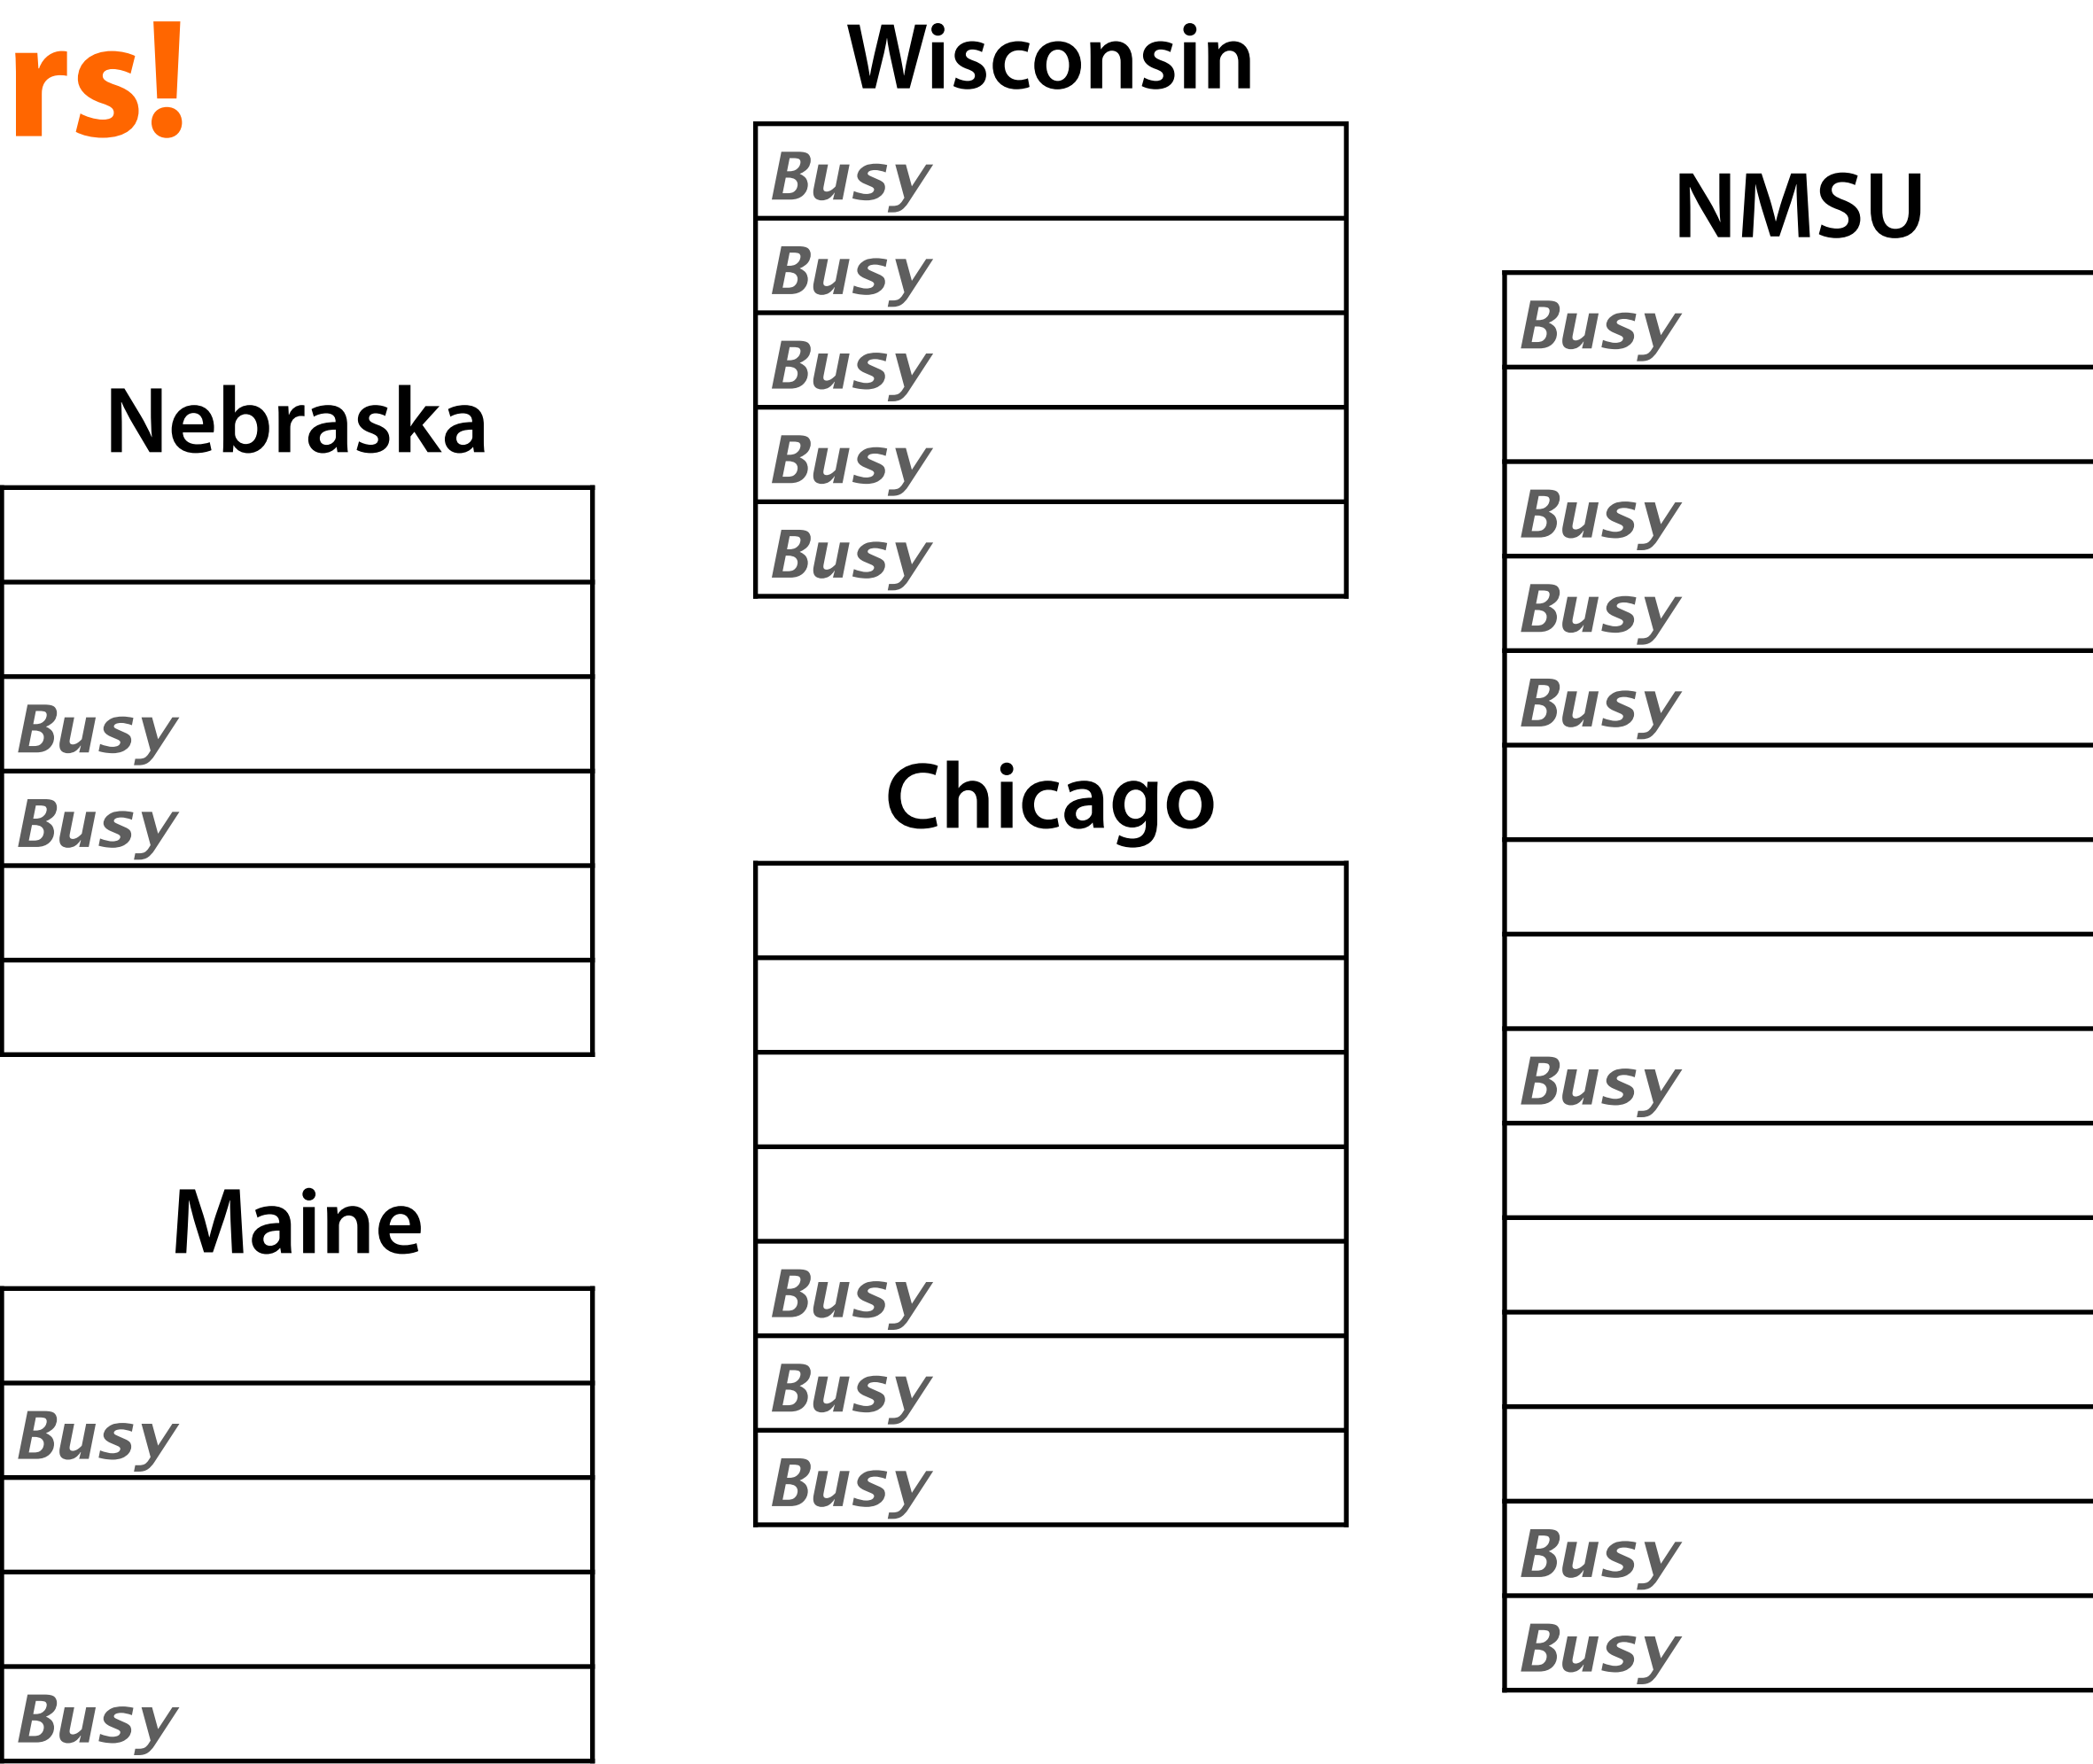
<i>Busy</i>
<i>Busy</i>
<i>Busy</i>
<i>Busy</i>
<i>Busy</i>



## 2. Add capacity contributors!

### Access Point

Job1.0  
Job1.1  
Job1.2  
Job1.3  
...  
Job1.1999







## 3. Request capacity (method #1)

Start Execution Points on clusters

### Access Point

Job1.0  
Job1.1  
Job1.2  
Job1.3  
...  
Job1.1999

#### Nebraska

OSPool EP NU1
OSPool EP NU2
Busy
Busy

#### Maine

OSPool EP ME2
Busy
OSPool EP ME1
OSPool EP ME3
Busy

#### Wisconsin

Busy
Busy
Busy
Busy
Busy

#### Chicago

OSPool EP UC2
OSPool EP UC1
Busy
Busy
Busy

#### NMSU

Busy
OSPool EP NM1
Busy
Busy
Busy
OSPool EP NM4
OSPool EP NM2
OSPool EP NM3
Busy
Busy
Busy



## 4. EPs add capacity to Pool

(I am not explaining how yet)

**AP**

Job1.0  
Job1.1  
Job1.2  
Job1.3  
...  
Job1.1999

**Pool**

OSPoo1 EP NU1	idle
OSPoo1 EP NU2	idle
OSPoo1 EP ME1	idle
OSPoo1 EP ME2	idle
OSPoo1 EP ME3	idle
OSPoo1 EP UC1	idle
OSPoo1 EP UC2	idle
OSPoo1 EP NM1	idle
OSPoo1 EP NM2	idle
OSPoo1 EP NM3	idle
OSPoo1 EP NM4	idle

**Nebraska**

OSPoo1 EP NU1
OSPoo1 EP NU2
Busy
Busy

**Maine**

OSPoo1 EP ME2
Busy
OSPoo1 EP ME1
OSPoo1 EP ME3
Busy

**Wisconsin**

Busy
Busy
Busy
Busy
Busy

**Chicago**

OSPoo1 EP UC2
OSPoo1 EP UC1
Busy
Busy
Busy

**NMSU**

Busy
OSPoo1 EP NM1
Busy
Busy
Busy
OSPoo1 EP NM4
OSPoo1 EP NM2
OSPoo1 EP NM3
Busy
Busy
Busy




## 5. Run jobs

HTCondor with AP & Pool

### AP

Job1.0  
Job1.1  
Job1.2  
Job1.3  
...  
Job1.1999



### Pool

OSPoo1	EP	NU1	Job1.4
OSPoo1	EP	NU2	idle
OSPoo1	EP	ME1	Job1.0
OSPoo1	EP	ME2	Job1.3
OSPoo1	EP	ME3	idle
OSPoo1	EP	UC1	Job1.2
OSPoo1	EP	UC2	Job1.6
OSPoo1	EP	NM1	Job1.8
OSPoo1	EP	NM2	Job1.12
OSPoo1	EP	NM3	Job1.10
OSPoo1	EP	NM4	idle

### Nebraska

NU1 > Job1.4
NU2 > idle
Busy
Busy

### Maine

ME2 > Job1.3
Busy
ME1 > Job1.0
ME3 > idle
Busy

### Wisconsin

Busy
Busy
Busy
Busy
Busy

### Chicago

UC2 > Job1.6
UC1 > Job1.2
Busy
Busy
Busy

### NMSU

Busy
NM1 > Job1.8
Busy
Busy
Busy
NM4 > idle
NM2 > Job1.12
NM3 > Job1.10
Busy
Busy
Busy



## 6. Get resources (method #2)

### Direct contributions

**AP**

Job1.0  
Job1.1  
Job1.2  
Job1.3  
...  
Job1.1999

Pool			
OSPoo1	EP	NU1	Job1.4
OSPoo1	EP	NU2	idle
OSPoo1	EP	ME1	Job1.0
OSPoo1	EP	ME2	Job1.3
OSPoo1	EP	ME3	idle
OSPoo1	EP	UC1	Job1.2
OSPoo1	EP	UC2	Job1.6
OSPoo1	EP	NM1	Job1.8
OSPoo1	EP	NM2	Job1.12
OSPoo1	EP	NM3	Job1.10
OSPoo1	EP	NM4	idle
NM	Contrb	EP1	idle

Nebraska	
NU1	> Job1.4
NU2	> idle
Busy	
Busy	

Maine	
ME2	> Job1.3
Busy	
ME1	> Job1.0
ME3	> idle
Busy	

Wisconsin	
Busy	
Busy	
Busy	
Busy	
Busy	

Chicago	
UC2	> Job1.6
UC1	> Job1.2
Busy	
Busy	
Busy	

NMSU	
Busy	
NM1	> Job1.8
Busy	
Busy	
Busy	
NM4	> idle
NM2	> Job1.12
NM3	> Job1.10
Busy	
NM Contrb EP1	
Busy	
Busy	

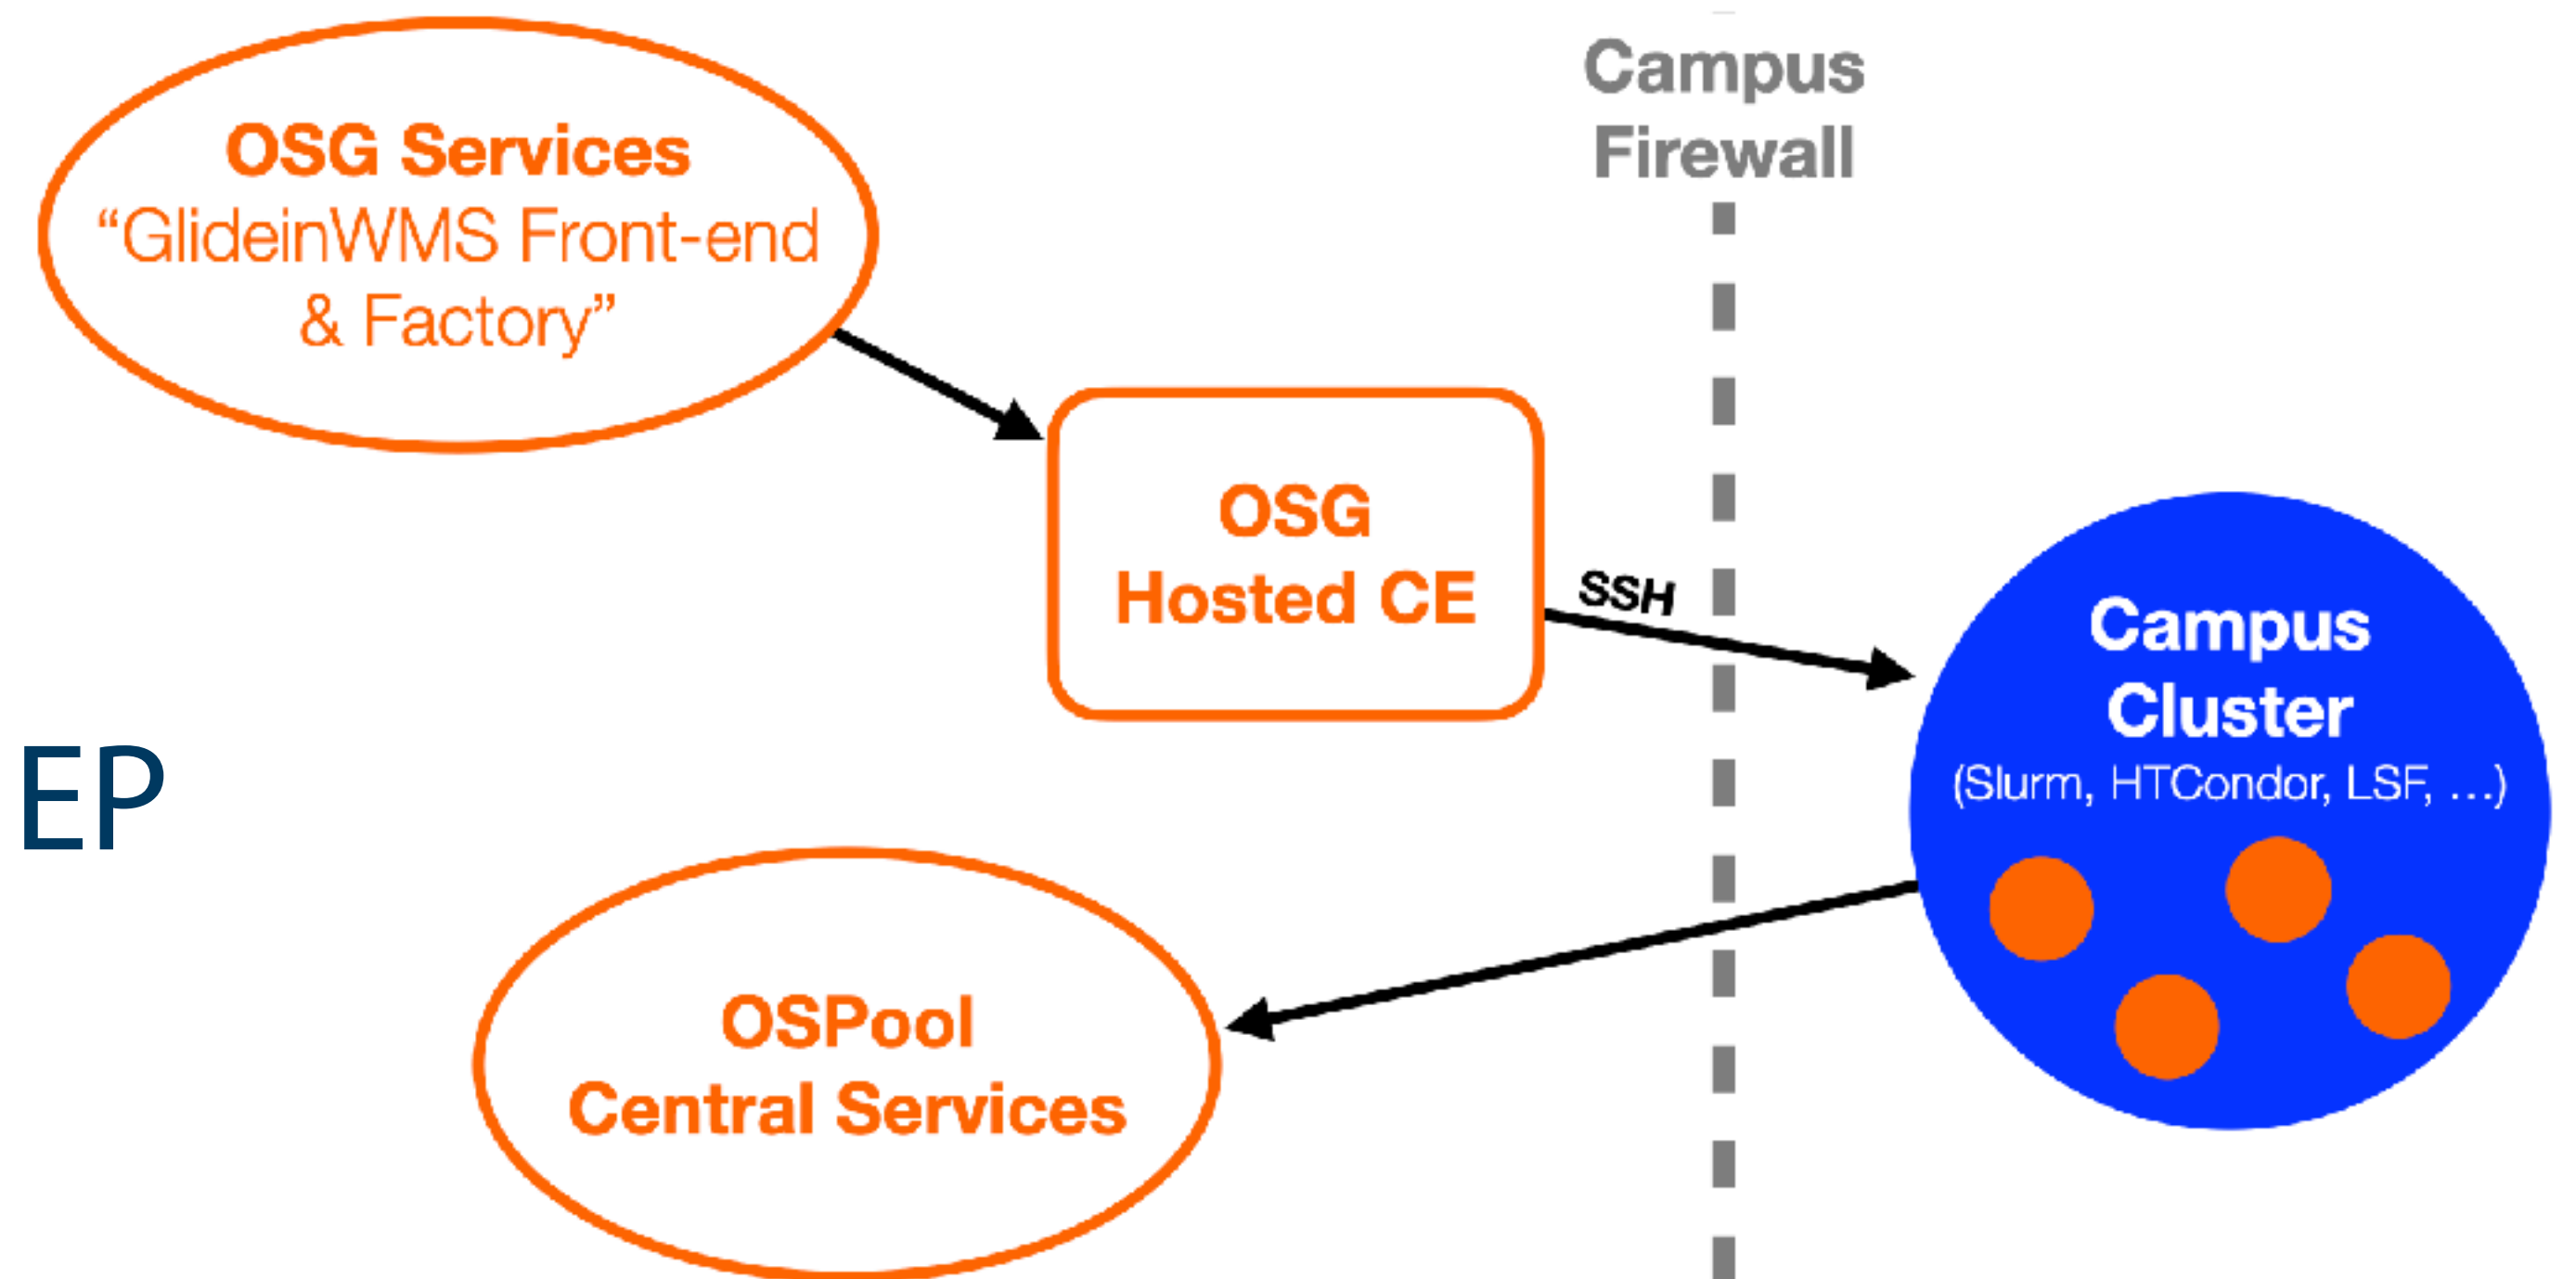


## For a batch scheduler:

- CE requests capacity (as jobs) based on demand
- Scheduler may run req.s
- Our SW creates *Execution Point* & adds to OSPool
- OSPool workloads run on EP

## Using containers:

- Admin starts containers
- SW makes EP, etc. (same)







# OSPool Contributors (*United States*)









# OSPool Usage



## Jobs



## Hours





# Using OSPool



# OSPool Is HTCondor



- OSPool is an HTCondor pool:  
You have condor\_q, condor\_submit, DAGMan, etc.
- OSPool bonus features!
  - More capacity (usually) than a typical local system
  - Some storage on Access Point (Data lecture, Wed.)
  - Some special resources, like GPUs (AI/ML topic, Thu.)
- How does OSPool differ from local one? *Variety*



# Varied Hardware



- Wider variety of CPUs (type, speed), memory, ...
- Request what you need in submit files  
(**request\_cpus**, **request\_memory**, **request\_disk**)
- Some specific hardware may be specified; search for documentation or contact us
  - Often in submit-file **requirements** expression
  - Example: GPU needs (ask us!)





- Varied Operating Systems
  - All Linux, *mostly* recent, but lots of variation
  - Changes to CentOS 8 => new variants (e.g., Alma)
  - Software on the Access Point **probably won't exist** on Execution Points! (e.g., specific Python version)
- Your software
  - Never assume your software is on Execution Points
  - The Software lecture (later today) is on this topic!



# Varied Access to Data



- No shared filesystem
  - Unlike some local clusters with shared filesystems
  - Thus, files must be transferred to Execution Points
- There are many ways to handle data
  - Data lecture is Wednesday morning



- Individual sites/clusters have their own policies
  - Example: Whether to kill jobs that exceed memory
  - Example: Maximum run-time of a job (or its glidein)
- If possible, set requirements for what you need
  - But this does not help with, e.g., maximum run-time
- Generally, try to make “OSG-sized” jobs (see next)



# What Makes a Good OSG Job?



	Ideal Jobs! (up to 10,000 cores across Jobs, per user!)	Still Very Advantageous!	Less-so, but maybe
<b>Cores</b> (GPUs)	<b>1</b> (1; non-specific type)	<b>&lt;8</b> (1; specific GPU type)	<b>&gt;8 (or MPI)</b> (multiple)
<b>Walltime</b>	<b>&lt;10 hrs*</b> *or checkpointable	<b>&lt;20 hrs*</b> *or checkpointable	<b>&gt;20 hrs</b>
<b>RAM</b>	<b>&lt;few GB</b>	<b>&lt;10s GB</b>	<b>&gt;10s GB</b>
<b>Input</b>	<b>&lt;500 MB</b>	<b>&lt;10 GB</b>	<b>&gt;10 GB</b>
<b>Output</b>	<b>&lt;1 GB</b>	<b>&lt;10 GB</b>	<b>&gt;10 GB</b>
<b>Software</b>	<i>'portable' (pre-compiled binaries, transferable, containerizable, etc.)</i>	<i>most other than</i> →	<i>Licensed software; non-Linux</i>



# More OSG Tips – Security



- Computer security is hard — read the headlines!
- OSG does its best, but no system is perfect
- Some suggestions:
  - Use strong, distinct passwords for each account
  - Do not share your account
  - Avoid world-writable directories and files
  - Avoid sensitive software and data (no HIPAA!)
  - Do not try to work around security barriers;  
contact us to help meet your goals in a safe way



# Acknowledgements





# ***You Can Acknowledge OSG!***



If you publish or present results that benefitted from using OSG services, please acknowledge us!

<https://osg-htc.org/acknowledging>



# Acknowledgements



- OSG team, especially Christina Koch; in past years: Brian Lin, Mats Rynge, and Jason Patton
- This work was supported by the National Science Foundation under Cooperative Agreement OAC-2030508



# A Few Suggestions



- **Exercises**
  - Today, some exercises will specify less, so try to use what you learned yesterday — first, from memory, if possible, then look things up
  - Use Slack! There are staff online who can help, too
- **Consultations**
  - Please consider signing up for a consultation, if you haven't already
  - If the slots fill up, we will try to make more!



# Demo, Part II